

STA 141C R Code

Jenny Xu

2025-05-26

```
library(tidyverse)

## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr   1.5.0
## v ggplot2     3.5.2      v tibble    3.2.1
## v lubridate  1.9.3      v tidyr     1.3.0
## v purrr      1.0.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors

library(ggplot2)
library(pheatmap)
library(caret)

## Loading required package: lattice
##
## Attaching package: 'caret'
##
## The following object is masked from 'package:purrr':
##
##     lift

library(glmnet)

## Loading required package: Matrix
##
## Attaching package: 'Matrix'
##
## The following objects are masked from 'package:tidyr':
##
##     expand, pack, unpack
##
## Loaded glmnet 4.1-8

library(e1071)
library(rpart)
library(rpart.plot)
library(randomForest)

## randomForest 4.7-1.1
## Type rfNews() to see new features/changes/bug fixes.
##
```

```

## Attaching package: 'randomForest'
##
## The following object is masked from 'package:dplyr':
##
##     combine
##
## The following object is masked from 'package:ggplot2':
##
##     margin
library(class)
library(MLmetrics)

##
## Attaching package: 'MLmetrics'
##
## The following objects are masked from 'package:caret':
##
##     MAE, RMSE
##
## The following object is masked from 'package:base':
##
##     Recall
library(yardstick)

##
## Attaching package: 'yardstick'
##
## The following objects are masked from 'package:caret':
##
##     precision, recall, sensitivity, specificity
##
## The following object is masked from 'package:readr':
##
##     spec
library(dplyr)
library(pROC)

## Type 'citation("pROC")' for a citation.
##
## Attaching package: 'pROC'
##
## The following objects are masked from 'package:stats':
##
##     cov, smooth, var
library(scales)

##
## Attaching package: 'scales'
##
## The following object is masked from 'package:purrr':
##
##     discard
##

```

```
## The following object is masked from 'package:readr':
##
##     col_factor
```

Exploratory Data Analysis

```
# Exploratory Data Analysis
# Read data
data = read.csv("Breast_GSE45827.csv")

# View shape and basic structure of the dataset
dim(data)

## [1] 151 54677
head(data[,1:10])

##  samples  type X1007_s_at X1053_at  X117_at  X121_at X1255_g_at X1294_at
## 1      84 basal  9.850040 8.097927 6.424728 7.353027  3.029122 6.880079
## 2      85 basal  9.861357 8.212222 7.062593 7.685578  3.149468 7.542283
## 3      87 basal 10.103478 8.936137 5.735970 7.687822  3.125931 6.562369
## 4      90 basal  9.756875 7.357148 6.479183 6.986624  3.181638 7.802344
## 5      91 basal  9.408330 7.746404 6.693980 7.333426  3.169923 7.610457
## 6      92 basal  7.505488 8.802820 6.235074 7.202227  2.987976 7.985281
##  X1316_at X1320_at
## 1 4.963740 4.408328
## 2 5.129607 4.584418
## 3 4.813449 4.425195
## 4 5.490982 4.567956
## 5 5.372469 4.424426
## 6 5.413368 4.465616

# Check the unique values and levels of outcome variable
table(data$type)

##
##      basal cell_line      HER luminal_A luminal_B      normal
##       41       14       30       29       30       7

unique(data$type)

## [1] "basal"      "HER"         "cell_line"  "normal"     "luminal_A"  "luminal_B"

# Check missing values -- there's no missing values in our dataset
sum(is.na(data))

## [1] 0
```

Dataset Composition

```
# Dataset Composition
# Count per class
type_counts <- data %>%
  count(type) %>%
  mutate(Percentage = n / sum(n) * 100)
```

```

# Bar plot of class counts
bar_plot <- ggplot(type_counts, aes(x = type, y = n, fill = type)) +
  geom_bar(stat = "identity") +
  labs(title = "Sample Count by Class", x = "Cancer Subtype", y = "Count") +
  theme_minimal() +
  theme(legend.position = "none")

# Pie chart of class proportions
pie_plot <- ggplot(type_counts, aes(x = "", y = Percentage, fill = type)) +
  geom_bar(stat = "identity", width = 1) +
  coord_polar("y") +
  labs(title = "Class Distribution (Pie Chart)") +
  theme_void() +
  geom_text(aes(label = paste0(round(Percentage, 1), "%")),
            position = position_stack(vjust = 0.5), color = "white", size = 4)

# Combine the two plots using gridExtra
library(gridExtra)


##
## Attaching package: 'gridExtra'

## The following object is masked from 'package:randomForest':
##
##      combine

## The following object is masked from 'package:dplyr':
##
##      combine

grid.arrange(bar_plot, pie_plot, ncol = 2)

```



Descriptive Statistics by Gene

```
# Descriptive Statistics by Gene
# remove first 2 columns (samples and types)
gene_data = data[, -(1:2)]

gene_summary = tibble(
  Gene = colnames(gene_data),
  Mean = colMeans(gene_data, na.rm = TRUE),
  Median = apply(gene_data, 2, median, na.rm = TRUE),
  SD = apply(gene_data, 2, sd, na.rm = TRUE),
  Min = apply(gene_data, 2, min, na.rm = TRUE),
  Max = apply(gene_data, 2, max, na.rm = TRUE),
)

head(gene_summary)
```

```
## # A tibble: 6 x 6
##   Gene      Mean Median    SD   Min   Max
##   <chr>    <dbl>  <dbl> <dbl> <dbl> <dbl>
## 1 X1007_s_at 10.3   10.4  0.613  7.51 11.7
## 2 X1053_at   7.63   7.53  0.706  5.86  9.63
## 3 X117_at    6.22   6.24  0.645  4.76  8.36
## 4 X121_at    7.34   7.33  0.331  6.63  8.37
## 5 X1255_g_at 3.19   3.19  0.159  2.76  3.61
## 6 X1294_at   7.31   7.42  0.642  5.46  8.57
```

Descriptive Statistics by Sample

```
# Compute summary statistics per row (i.e., per sample)
sample_summary <- tibble(
  SampleID = data$samples,
  Type     = data$type,
  Mean     = apply(gene_data, 1, mean, na.rm = TRUE),
  Median   = apply(gene_data, 1, median, na.rm = TRUE),
  SD       = apply(gene_data, 1, sd, na.rm = TRUE),
  Min      = apply(gene_data, 1, min, na.rm = TRUE),
  Max      = apply(gene_data, 1, max, na.rm = TRUE)
)

# Preview the result
head(sample_summary)
```

```
## # A tibble: 6 x 7
##   SampleID Type   Mean Median   SD   Min   Max
##   <int> <chr> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1     84 basal  5.66  5.17  2.16  2.44  14.8
## 2     85 basal  5.67  5.19  2.18  2.51  14.7
## 3     87 basal  5.65  5.15  2.17  2.34  14.7
## 4     90 basal  5.66  5.19  2.16  2.41  14.7
## 5     91 basal  5.66  5.19  2.17  2.46  14.7
## 6     92 basal  5.67  5.20  2.16  2.38  14.7
```

Descriptive Statistics by Breast Cancer Subtype

```
# Pivot to long format
data_long = data %>%
  pivot_longer(cols = -(samples:type), names_to = "Gene", values_to = "Expression")

# Compute stats by class and gene
grouped_summary = data_long %>%
  group_by(type, Gene) %>%
  summarise(
    Mean = mean(Expression, na.rm = TRUE),
    Median = median(Expression, na.rm = TRUE),
    SD = sd(Expression, na.rm = TRUE),
    .groups = "drop"
  )

# Preview
head(grouped_summary)
```

```
## # A tibble: 6 x 5
##   type Gene           Mean Median   SD
##   <chr> <chr>         <dbl> <dbl> <dbl>
## 1 HER  AFFX.BioB.3_at  7.70  7.62 0.257
## 2 HER  AFFX.BioB.5_at  7.85  7.72 0.334
## 3 HER  AFFX.BioB.M_at  8.36  8.32 0.265
## 4 HER  AFFX.BioC.3_at  9.49  9.43 0.279
## 5 HER  AFFX.BioC.5_at  9.18  9.12 0.233
## 6 HER  AFFX.BioDn.3_at 11.8  11.8 0.217
```

Overall Descriptive Statistics

```
# Overall Descriptive Statistics
# Flatten to a vector
all_values = unlist(gene_data, use.names = FALSE)

overall_summary = tibble(
  # Gene = colnames(gene_data),
  Mean = mean(all_values, na.rm = TRUE),
  Median = median(all_values, na.rm = TRUE),
  SD = sd(all_values, na.rm = TRUE),
  Min = min(all_values, na.rm = TRUE),
  Max = max(all_values, na.rm = TRUE),
)

head(overall_summary)

## # A tibble: 1 x 5
##   Mean Median    SD   Min   Max
##   <dbl> <dbl> <dbl> <dbl> <dbl>
## 1  5.65  5.17  2.13  2.17  15.0

print(paste("Mean:", overall_summary$Mean))

## [1] "Mean: 5.65347002674015"

print(paste("Median:", overall_summary$Median))

## [1] "Median: 5.17401561951104"

print(paste("Standard Deviation:", overall_summary$SD))

## [1] "Standard Deviation: 2.12764795656006"

print(paste("Min:", overall_summary$Min))

## [1] "Min: 2.17109972663613"

print(paste("Max:", overall_summary$Max))

## [1] "Max: 14.9701002149474"
```

Max/Min of Mean Gene Expression

```
# Sample with highest mean expression
max_mean_sample <- sample_summary %>%
  filter(Mean == max(Mean, na.rm = TRUE))

# Sample with lowest mean expression
min_mean_sample <- sample_summary %>%
  filter(Mean == min(Mean, na.rm = TRUE))

# Print results
cat("Sample with MAX mean expression:\n")

## Sample with MAX mean expression:
```

```
print(max_mean_sample)

## # A tibble: 1 x 7
##   SampleID Type      Mean Median    SD   Min   Max
##   <int> <chr>    <dbl>  <dbl> <dbl> <dbl> <dbl>
## 1      183 luminal_B  5.68   5.18  2.19  2.45  14.7

cat("\nSample with MIN mean expression:\n")
```

```
##
## Sample with MIN mean expression:
```

```
print(min_mean_sample)

## # A tibble: 1 x 7
##   SampleID Type      Mean Median    SD   Min   Max
##   <int> <chr>    <dbl>  <dbl> <dbl> <dbl> <dbl>
## 1      220 luminal_A  5.59   5.15  2.05  2.44  14.7
```

Max/Min of Absolute Gene Expression

```
# Sample with highest expression
max_sample <- sample_summary %>%
  filter(Max == max(Max, na.rm = TRUE))

# Sample with lowest expression
min_sample <- sample_summary %>%
  filter(Min == min(Min, na.rm = TRUE))

# Print results
cat("Sample with MAX mean expression:\n")
```

```
## Sample with MAX mean expression:
```

```
print(min_sample)

## # A tibble: 1 x 7
##   SampleID Type      Mean Median    SD   Min   Max
##   <int> <chr>    <dbl>  <dbl> <dbl> <dbl> <dbl>
## 1      147 basal  5.64   5.19  2.08  2.17  14.8

cat("\nSample with MIN mean expression:\n")
```

```
##
## Sample with MIN mean expression:
```

```
print(max_sample)

## # A tibble: 1 x 7
##   SampleID Type      Mean Median    SD   Min   Max
##   <int> <chr>    <dbl>  <dbl> <dbl> <dbl> <dbl>
## 1      151 HER    5.66   5.17  2.16  2.37  15.0
```

Frequency distribution of Gene Expression - Why we need normalization/scaling

```
# Distribution of Gene Expression - Why we need normalization/scaling
all_values = unlist(gene_data, use.names = FALSE)
```



```

expression_df = tibble(Expression = all_values)

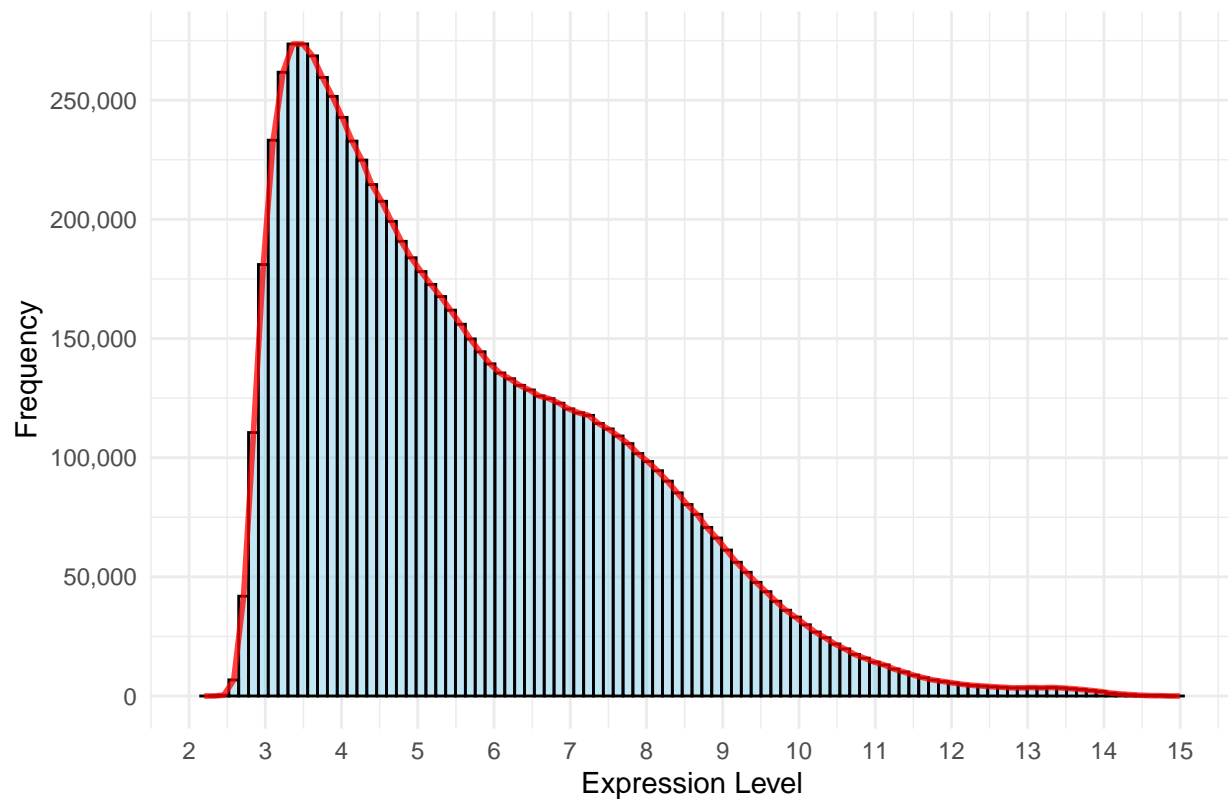
ggplot(expression_df, aes(x = Expression)) +
  geom_histogram(bins = 100, fill = "skyblue", color = "black", alpha = 0.5) +
  geom_line(
    stat = "bin",
    bins = 100,
    aes(y = ..count..),
    color = "red",
    size = 1,
    alpha = 0.75
  ) +
  labs(title = "Distribution of All Gene Expression Values",
       x = "Expression Level", y = "Frequency") +
  scale_x_continuous(breaks = seq(0, 20, by = 1)) +
  scale_y_continuous(
    labels = scales::label_comma(),
    breaks = seq(0, 500000, by = 50000)
  ) +
  theme_minimal()

## Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.
## i Please use `linewidth` instead.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.

## Warning: The dot-dot notation (`..count..`) was deprecated in ggplot2 3.4.0.
## i Please use `after_stat(count)` instead.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.

```

Distribution of All Gene Expression Values



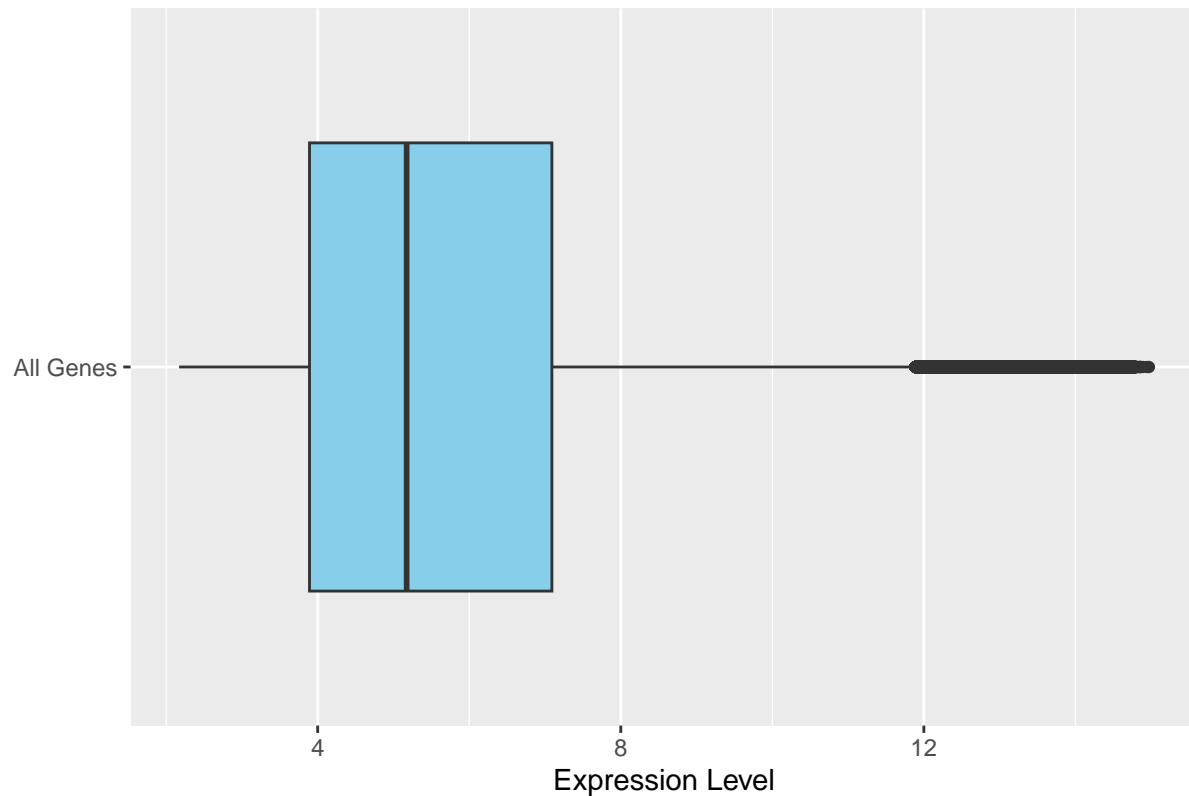
Boxplot of Gene Expression - Why we need normalization/scaling

```
# Boxplot of Gene Expression - Why we need normalization/scaling
# Prepare data: expression as y-axis
expression_df = tibble(Expression = unlist(data[ , -(1:2)]), use.names = FALSE))

# Add dummy variable for plotting
expression_df$Group = "All Genes"

ggplot(expression_df, aes(y = Group, x = Expression)) +
  geom_boxplot(fill = "skyblue") +
  labs(title = "Overall Distribution of Gene Expression Values",
       y = "", x = "Expression Level")
```

Overall Distribution of Gene Expression Values

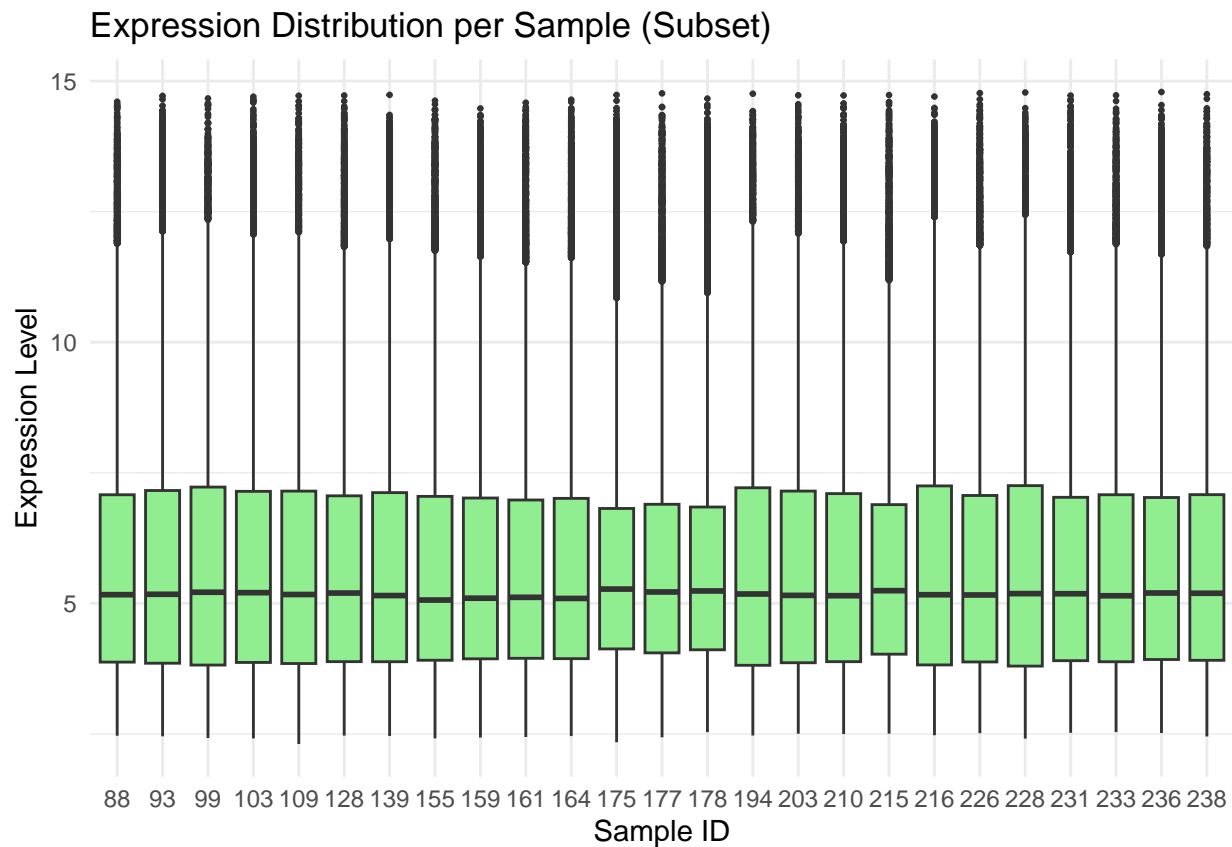


Boxplot of Gene Expression by Sample

```
set.seed(123)
# Boxplot for a few random samples to keep it readable
sample_subset = sample(unique(data_long$samples), 25)
plot_data = data_long %>% filter(samples %in% sample_subset)

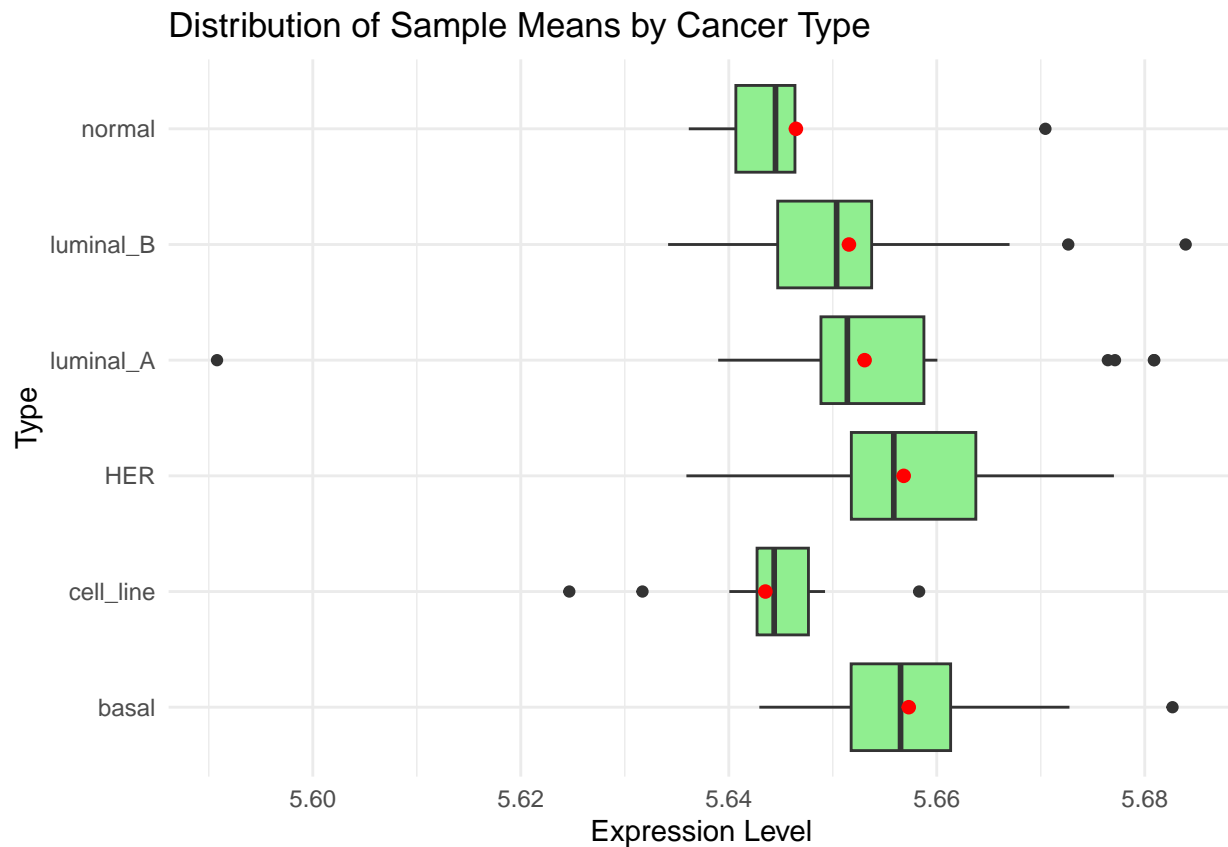
# plot_data = data_long %>% filter(samples %in% data_long$samples)

ggplot(plot_data, aes(x = factor(samples), y = Expression)) +
  geom_boxplot(outlier.size = 0.5, fill = "lightgreen") +
  labs(title = "Expression Distribution per Sample (Subset)",
       x = "Sample ID", y = "Expression Level") +
  theme_minimal()
```



Boxplot of Gene Expression by Cancer Type

```
ggplot(sample_summary, aes(y = Type, x = Mean)) +
  geom_boxplot(fill = "lightgreen") +
  stat_summary(fun = mean, geom = "point", shape = 20, size = 3, color = "red") +
  labs(title = "Distribution of Sample Means by Cancer Type",
       x = "Expression Level") +
  theme_minimal()
```




Top 10 Variable Genes - Do their gene expression differ by cancer subtype?

```
gene.variances = apply(gene_data, 2, var)
top.genes = sort(gene.variances, decreasing = TRUE)
top.10.genes.name = names(top.genes)[1:10]

# Create a long-format dataframe for plotting
plot_data = lapply(top.10.genes.name, function(gene) {
  data.frame(
    Gene = gene,
    Expression = gene_data[[gene]],
    Subtype = data$type
  )
}) %>% bind_rows()

# Create combined boxplot
ggplot(plot_data, aes(x = Subtype, y = Expression, fill = Subtype)) +
  geom_boxplot(outlier.size = 0.5) +
  facet_wrap(~ Gene, scales = "free_y") +
  labs(title = "Top 10 Most Variable Genes: Expression by Subtype",
       x = "Cancer Subtype",
       y = "Gene Expression") +
  theme_bw() +
  theme(legend.position = "none")
```

Top 10 Most Variable Genes: Expression by Subtype



```
# Anova table
anova_results <- lapply(top.10.genes.name, function(gene) {
  model <- aov(gene_data[[gene]] ~ data$type)
  summary_df <- summary(model)[[1]]
  data.frame(
    Gene = gene,
    F_value = summary_df[["F value"]][1],
    p_value = summary_df[["Pr(>F)"]][1]
  )
}) %>% bind_rows()

# Print to console
print(anova_results)
```

```
##           Gene    F_value    p_value
## 1  X206378_at  21.58804 4.125586e-16
## 2  X228241_at 134.88529 1.053774e-52
## 3  X205916_at  12.04787 9.138208e-10
## 4  X237339_at  56.74219 1.909145e-32
## 5  X1553602_at  18.87886 1.989393e-14
## 6  X209173_at  86.16018 1.178749e-41
## 7  X214451_at  18.90565 1.912701e-14
## 8  X205009_at  44.16762 1.643301e-27
## 9  X205044_at  35.16204 1.948368e-23
## 10 X202917_s_at 29.47704 1.438317e-20
```

```
# for (gene in top.10.genes.name){
#   # New data frame
```

```

# data.plot = data.frame(
#   expression = gene_data[[gene]],
#   subtype = data$type
# )
#
# # Boxplot of how gene expression differ by subtype for highly variable gene
# print(
#   ggplot(data.plot, aes(x = subtype, y = expression, fill = subtype)) +
#   geom_boxplot() +
#   labs(title = paste("Expression of", gene, "by Subtype"),
#        x = "Subtype",
#        y = "Expression Level")
# )
# # One-way ANOVA for one top gene
# print(summary(aov(expression~subtype, data = data.plot)))
# }


```

PCA Plot - Sample Cluster by Cancer Subtype

```

# PCA plot -- samples cluster by cancer subtype
# Standardize all genes (mean 0, sd 1)
scaled.data = scale(gene_data)
# PCA
pca.result = prcomp(scaled.data)
pca.plot = data.frame(
  PC1 = pca.result$x[,1],
  PC2 = pca.result$x[,2],
  Subtype = data$type
)
ggplot(pca.plot, aes(x = PC1, y = PC2, color = Subtype)) +
  geom_point(size = 2) +
  labs(title = "PCA of Gene Expression (Scaled)", x = "PC1", y = "PC2")

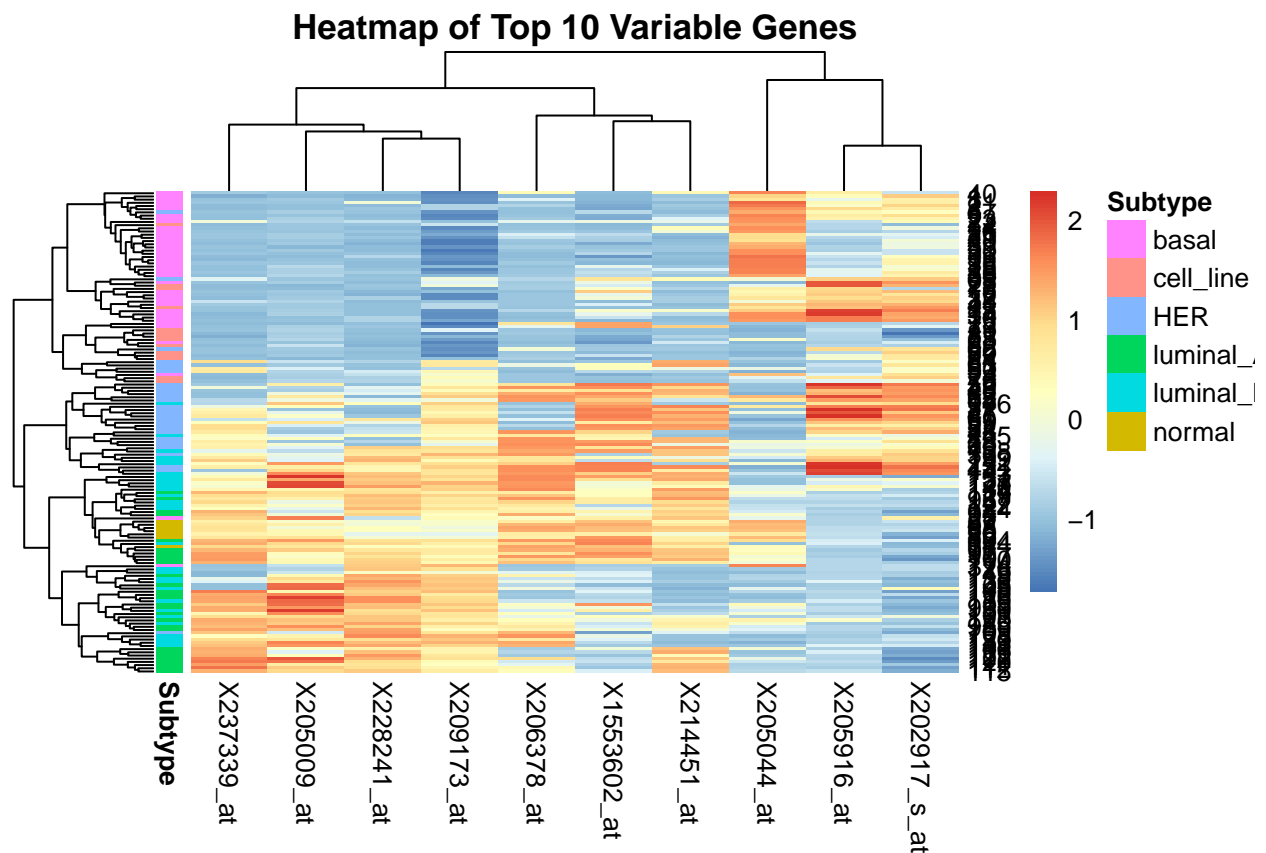
```



Heatmap - Sample Cluster by Subtype

```
# Heatmap of 10 top variable genes
top.10.expression = gene_data[,top.10.genes.name]
top.10.scaled = scale(top.10.expression)
rownames(top.10.scaled) = rownames(data)
row.annotation = data.frame(Subtype = data$type)
rownames(row.annotation) = rownames(top.10.scaled)

pheatmap(top.10.scaled,
  annotation_row = row.annotation,
  cluster_rows = TRUE,
  cluster_cols = TRUE,
  main = "Heatmap of Top 10 Variable Genes")
```

Methodology

Data Pre-processing

```
# Methodology (Pre-processing steps)
# Step 1: Split into Training and Test sets
set.seed(141)
idx <- createDataPartition(data$type, p = 0.80, list = FALSE)
train.data <- data[idx, ]
test.data <- data[-idx, ]

table(test.data$type)
```

```
##
##      basal cell_line      HER luminal_A luminal_B      normal
##          8         2         6         5         6         1
```

```
# Step 2: Extract gene expression matrix and subtype labels
```

```
x.train.raw = train.data[,-(1:2)]
y.train = train.data$type
x.test.raw = test.data[,-(1:2)]
y.test = test.data$type
```

```
# Step 3: Scale training data only and test data using training parameters
```

```
scaled.train = scale(x.train.raw)
train.center = attr(scaled.train, "scaled:center")
```

```

train.scale = attr(scaled.train, "scaled:scale")
scaled.test = scale(x.test.raw,
                    center = train.center,
                    scale = train.scale)

# Step 4: Handle class imbalance by upsampling scaled training data
y.train = factor(train.data$type)

train.upsampled = upSample(
  x = as.data.frame(scaled.train),
  y = y.train,
  yname = "target"
)

x.train.resampled = train.upsampled[, -ncol(train.upsampled)]
y.train.resampled = train.upsampled$target

```

80% Training vs. 20% Test Split (Before Upsampling)

```

# Before upsampling, 80% Train set vs. 20% Test set
# build a tiny data frame
n_train = length(y.train)
n_test = length(y.test)
df <- data.frame(
  Set = factor(c("Train", "Test"), levels = c("Train", "Test")),
  Count = c(n_train, n_test)
)

# plot
ggplot(df, aes(x = Set, y = Count, fill = Set)) +
  geom_col(width = 0.6) + # geom_col is the same as geom_bar(stat="identity")
  labs(title = "Train vs Test Set Sizes",
       x = NULL,
       y = "Number of samples") +
  theme_minimal() +
  theme(legend.position = "none")

```



Unbalanced vs. Balanced Class (After Upsampling)

```
# After upsampling, balanced class in training set
# train.upsampled$target is the factor of up-sampled labels
table(train.upsampled$target)
```

```
##
##      basal cell_line      HER luminal_A luminal_B      normal
##          33          33          33          33          33          33
```

```
# class counts in your up-sampled test set
table(y.train.resampled)
```

```
## y.train.resampled
##      basal cell_line      HER luminal_A luminal_B      normal
##          33          33          33          33          33          33
```

```
table(test.data$type)
```

```
##
##      basal cell_line      HER luminal_A luminal_B      normal
##          8          2          6          5          6          1
```

PCA

```
# Methodology (PCA)
# Step 1: Apply PCA to scaled training data
pca.model = prcomp(x.train.resampled, center = FALSE, scale. = FALSE)
```

```

# Step 2: Examine variance explained by each component
explained.variance = pca.model$sdev^2
proportion.variance = explained.variance/sum(explained.variance)
cumulative.variance = cumsum(proportion.variance)

# Step 3: Find number of PCs that explain more than 95% of variance
num.pc.95 = which(cumulative.variance >= 0.95)[1]


# Step 4: Create PCA-transformed training data
pca.train = as.data.frame(pca.model$x[,1:num.pc.95])
pca.train$Subtype = y.train.resampled

# Step 5: Apply same PCA to test data
pca.test.raw = predict(pca.model, newdata = scaled.test)
pca.test = as.data.frame(pca.test.raw[,1:num.pc.95])

# Step 6: Scree plot: proportion of variance
plot(explained.variance[1:50], type = "b",
     xlab = "Principal Component",
     ylab = "Proportion of Variance Explained",
     main = "PCA Scree Plot (Top 50 PCs)")

```

PCA Scree Plot (Top 50 PCs)



```

# Step 7: Cumulative variance plot
screedata = data.frame(
  PC = 1:length(cumulative.variance),
  CumulativeVariance = cumulative.variance
)


ggplot(screedata, aes(x = PC, y = CumulativeVariance)) +

```

```

geom_line(color = "blue") +
geom_point() +
geom_hline(yintercept = 0.95, linetype = "dashed", color = "red") +
geom_vline(xintercept = num.pc.95, linetype = "dotted", color = "darkgreen") +
theme_minimal() +
labs(
  title = "Cumulative Variance Explained by PCA",
  x = "Number of Principal Components",
  y = "Cumulative Proportion of Variance Explained"
)

```



```

# What cumulative variance do those PCs explain?
cat("Cumulative variance at PC", num.pc.95, ":",
    round(cumulative.variance[num.pc.95], 4), "\n")

```

```
## Cumulative variance at PC 92 : 0.9515
```

Lasso

```

# Methodology (Lasso)
# Step 1: Prepare training and test data in matrix form
x.lasso.train = as.matrix(pca.train[, -ncol(pca.train)])
y.lasso.train = pca.train$Subtype


x.lasso.test = as.matrix(pca.test)

# Step 2: Fit Lasso Model

```

```
lasso.model = glmnet(x.lasso.train, y.lasso.train,
                    alpha = 1, family = "multinomial")

# Step 3: Use Cross-validation to find optimal lambda
set.seed(141)
cv.lasso = cv.glmnet(x.lasso.train, y.lasso.train,
                    alpha = 1, family = "multinomial")
plot(cv.lasso)
```



```
best.lambda = cv.lasso$lambda.min

# Step 4: Predict on test data
lasso.pred = predict(cv.lasso, newx = x.lasso.test,
                    s = best.lambda, type = "class")

# Step 5: Evaluate Model Performance
table(Predicted = lasso.pred, Actual = test.data$type)
```

```
##           Actual
## Predicted  basal cell_line HER luminal_A luminal_B normal
## basal      8           0  0           0           0       0
## cell_line   0           2  0           0           0       0
## HER         0           0  6           0           1       0
## luminal_A   0           0  0           5           1       0
## luminal_B   0           0  0           0           4       0
## normal      0           0  0           0           0       1
```

```
lasso.accuracy = mean(lasso.pred == test.data$type)
lasso.accuracy
```

```
## [1] 0.9285714
```

Which gene is most predictive for each breast cancer subtype using Lasso?

```
# Which gene is most predictive for each breast cancer subtype using Lasso?
# No scaling of raw data, just upsampling to avoid bias
# Step 1: Upsample raw data
set.seed(141)
train.raw.upsampled = upSample(
  x = x.train.raw,
  y = factor(y.train),
  yname = "target"
)

# Step 2: Prepare training and test data for Lasso
x.raw.lasso.train = as.matrix(train.raw.upsampled[, -ncol(train.raw.upsampled)])
y.raw.lasso.train = train.raw.upsampled$target

# Step 3: Fit lasso on raw gene data
lasso.raw.model = glmnet(x.raw.lasso.train, y.raw.lasso.train,
  alpha = 1, family = "multinomial")

# Step 4: Cross-validation for lambda
cv.raw.lasso = cv.glmnet(x.raw.lasso.train, y.raw.lasso.train,
  alpha = 1, family = "multinomial")
best.raw.lambda = cv.raw.lasso$lambda.min

# Step 5: Extract coefficients at best lambda
coef.raw = coef(cv.raw.lasso, s = best.raw.lambda)

# Step 6: Find top predictive genes per subtype
top.genes.by.subtype = list()

for (subtype in names(coef.raw)){
  coef.matrix = coef.raw[[subtype]]
  nonzero.idx = which(coef.matrix != 0)[-1]

  if (length(nonzero.idx) > 0){
    gene.name = rownames(coef.matrix)[nonzero.idx]
    abs.coef = abs(as.vector(coef.matrix[nonzero.idx]))
    names(abs.coef) = gene.name

    top.genes = sort(abs.coef, decreasing = TRUE)
    top.genes.by.subtype[[subtype]] = top.genes
  }else{
    top.genes.by.subtype[[subtype]] = NULL
  }
}

top.genes.by.subtype
```

```
## $basal
##   X1553315_at  X241044_x_at   X233730_at  X230226_s_at   X231374_at
##   0.644895206  0.486732596   0.460282225  0.329445630   0.279968846
##   X204580_at   X232440_at   X208358_s_at   X237246_at  X1557809_a_at
##   0.279843949  0.228475730   0.202975141   0.174167453   0.172691455
```

```

##      X205376_at X1564676_a_at      X220612_at      X214772_at      X238865_at
##      0.150992785      0.149102300      0.141474729      0.133711273      0.126005714
##      X223315_at X222457_s_at      X1554840_at      X233405_at      X227349_at
##      0.116169825      0.115616115      0.091852214      0.089881590      0.079107452
## X1568574_x_at      X205143_at X1553989_a_at      X226206_at      X208154_at
##      0.071489789      0.055545286      0.040303851      0.039905312      0.038286879
##      X205549_at X205029_s_at
##      0.007430235      0.001241784
##
## $cell_line
## X202878_s_at X201721_s_at      X211990_at X207365_x_at      X1569041_at      X214722_at
##      0.32624585      0.23698968      0.21831347      0.20950763      0.08842685      0.07470837
##      X235847_at X217757_at      X230332_at X210809_s_at      X219926_at
##      0.05808671      0.05013165      0.02278132      0.02077457      0.01994783
##
## $HER
##      X241884_at      X1565819_at      X210930_s_at      X236522_at      X1556923_at
##      1.43625764      1.07760441      0.99043397      0.40609026      0.21526034
##      X206793_at      X244162_at      X1557758_at X1552590_a_at      X1560556_a_at
##      0.16814060      0.14644032      0.14067245      0.13317818      0.11574505
##      X229306_at      X229194_at      X242275_at      X215802_at      X207284_s_at
##      0.06348842      0.05663565      0.05411003      0.05172053      0.04515138
##      X204915_s_at X1554712_a_at      X216917_s_at
##      0.04467974      0.03588161      0.03075034
##
## $luminal_A
##      X215014_at      X238625_at X201235_s_at      X229160_at      X243605_at      X206638_at
##      0.52484685      0.36448613      0.34387789      0.32474664      0.23392198      0.22426145
##      X223721_s_at X227742_at      X215856_at      X233977_at X231002_s_at      X227182_at
##      0.19448390      0.17139572      0.14332353      0.11030388      0.10366811      0.09468187
##      X209123_at      X218613_at X229110_at      X229461_x_at X202174_s_at      X205908_s_at
##      0.06338289      0.04710698      0.03878650      0.03428616      0.03411821      0.02759677
##      X243334_at      X242301_at
##      0.01911009      0.01581751
##
## $luminal_B
##      X205477_s_at      X228405_at      X1556654_at      X78047_s_at      X217724_at      X221836_s_at
##      0.9434511886      0.6873265248      0.5970686385      0.5945719647      0.5288954543      0.5191250060
##      X225090_at      X217351_at      X221811_at      X239278_at      X239612_at      X226727_at
##      0.4560141897      0.4453464140      0.4446205967      0.3335340055      0.3141100145      0.2382750585
##      X219401_at      X234046_at X234927_s_at      X211712_s_at      X219051_x_at      X203623_at
##      0.2155730559      0.1441536298      0.1132527216      0.1065677388      0.0692334567      0.0672823694
##      X214858_at X225350_s_at      X206107_at      X220148_at      X204378_at
##      0.0642107162      0.0486018096      0.0159176783      0.0056914940      0.0000139867
##
## $normal
##      X231598_x_at      X242641_at X1552509_a_at      X243689_s_at      X218872_at
##      0.67650609      0.57641822      0.49052294      0.28673783      0.22638400
##      X206093_x_at      X1561754_at      X211565_at
##      0.13479423      0.12250978      0.08008514

```


SVM

```
# Methodology (SVM)
# Step 1: Prepare training and test data from PCA transformed data
x.svm.train = as.matrix(pca.train[, -ncol(pca.train)])
y.svm.train = pca.train$Subtype

x.svm.test = as.matrix(pca.test)

# Step 2: Train SVM Model
set.seed(141)
svm.model = svm(
  x = x.svm.train,
  y = y.svm.train,
  kernel = "linear",
  probability = TRUE
)

# Step 3: Predict on Test data
svm.pred = predict(svm.model, newdata = x.svm.test)

# Step 4: Evaluate Accuracy
table(Predicted = svm.pred, Actual = test.data$type)
```

```
##           Actual
## Predicted  basal cell_line HER luminal_A luminal_B normal
## basal      8         0  0         0         0         0
## cell_line  0         2  0         0         0         0
## HER        0         0  6         0         0         0
## luminal_A  0         0  0         5         2         0
## luminal_B  0         0  0         0         4         0
## normal    0         0  0         0         0         1
```

```
svm.accuracy = mean(svm.pred == test.data$type)
svm.accuracy
```

```
## [1] 0.9285714
```

Decision Tree

```
# Methodology (Decision Tree)
# Step 1: Prepare training and test data from PCA transformed data
tree.train.data = pca.train
tree.test.data = pca.test

# Step 3: Fit tree model
tree.model = rpart(Subtype~., data = tree.train.data, method = "class")

# Step 4: Predict on test data
tree.pred = predict(tree.model, newdata = tree.test.data, type = "class")

# Step 5: Confusion matrix + accuracy
table(Predicted = tree.pred, Actual = test.data$type)
```

```
##           Actual
```

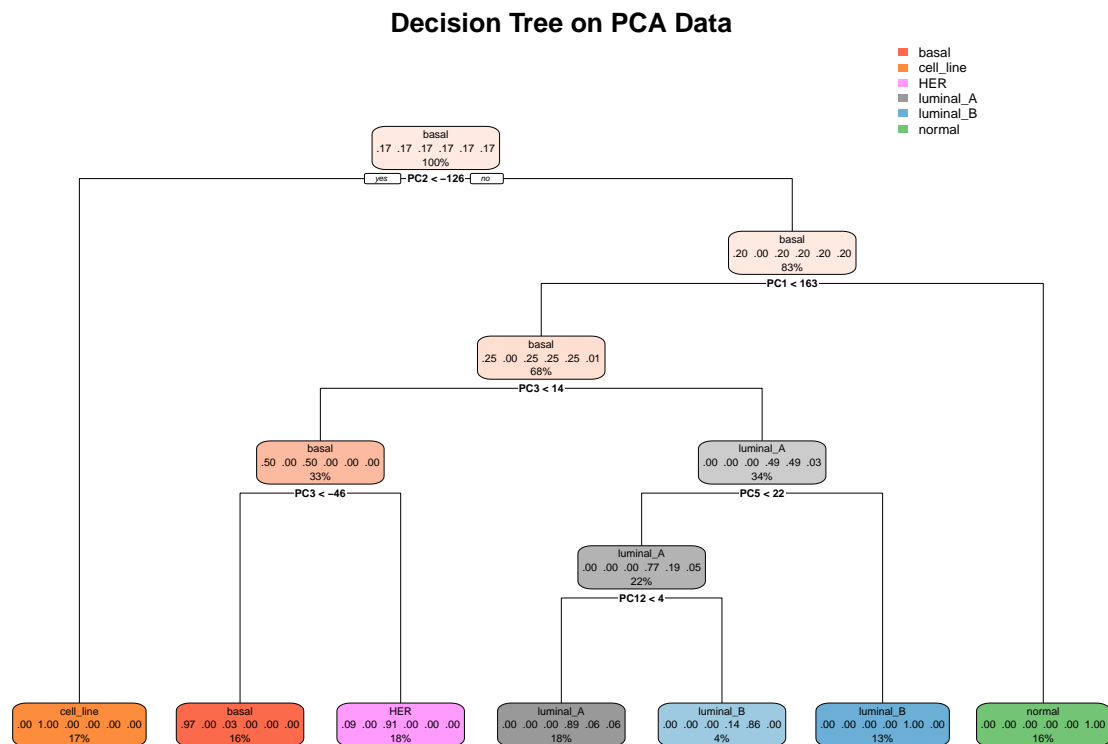
```
## Predicted   basal cell_line HER luminal_A luminal_B normal
## basal      8         0  0         0         0         0
## cell_line   0         2  0         0         0         0
## HER         0         0  6         0         0         0
## luminal_A   0         0  0         4         0         0
## luminal_B   0         0  0         1         6         0
## normal      0         0  0         0         0         1
```

```
tree.accruacy = mean(tree.pred == test.data$type)
tree.accruacy
```

```
## [1] 0.9642857
```

```
# Step 6: Visualize the tree
```

```
rpart.plot(tree.model, main = "Decision Tree on PCA Data")
```



Random Forest

```
# Methodology (Random Forest)
```

```
# Step 1: Prepare training and test data from PCA transformed data
```

```
rf.train.data = pca.train
```

```
rf.test.data = pca.test
```

```
# Step 3: Fit random forest model
```

```
set.seed(123)
```

```
rf.model = randomForest(Subtype~., data = rf.train.data, ntree = 500)
```

```
# Step 4: Predict on Test data
```

```
rf.pred = predict(rf.model, newdata = rf.test.data)
```

```
# Step 5: Confusion matrix + accuracy
```

```
table(Predicted = rf.pred, Actual = test.data$type)
```

```
##           Actual
## Predicted  basal cell_line HER luminal_A luminal_B normal
## basal      8         0  0         0         0         0
## cell_line   0         2  0         0         0         0
## HER         0         0  5         0         1         0
## luminal_A   0         0  0         4         1         0
## luminal_B   0         0  1         1         4         0
## normal     0         0  0         0         0         1
```

```
rf.accuracy = mean(rf.pred == test.data$type)
rf.accuracy
```

```
## [1] 0.8571429
```

KNN

```
# Methodology (KNN)
# Step 1: Prepare training and test data from PCA transformed data
x.knn.train = as.matrix(pca.train[, -ncol(pca.train)])
y.knn.train = pca.train$Subtype

x.knn.test = as.matrix(pca.test)

# Step 2: fit KNN model (manually choose k)
set.seed(141)
knn.pred = knn(train = x.knn.train, test = x.knn.test,
               cl = y.knn.train, k = 5)

# Step 3: Confusion matrix + accuracy
table(Predicted = knn.pred, Actual = test.data$type)
```

```
##           Actual
## Predicted  basal cell_line HER luminal_A luminal_B normal
## basal      7         0  1         0         0         0
## cell_line   0         2  0         0         0         0
## HER         1         0  2         0         0         0
## luminal_A   0         0  3         5         5         0
## luminal_B   0         0  0         0         1         0
## normal     0         0  0         0         0         1
```

```
knn.accuracy = mean(knn.pred == test.data$type)
knn.accuracy
```

```
## [1] 0.6428571
```

Main Results

Accuracy

```
# Main Results (Model Comparison using Caret package)
# Step 1: Extract features and target from PCA data
x.pca = pca.train[, -ncol(pca.train)]
```

```

y.pca = pca.train$Subtype

# Step 2: Define 5-fold CV and upsampling
ctrl <- trainControl(
  method = "cv",
  number = 5,
  sampling = "up"      # <-- this makes caret up-sample the minority classes in each fold
)

model.results.pca = list()

set.seed(141)
model.results.pca[["Random Forest"]] <- train(
  x = x.pca,
  y = y.pca,
  method = "rf",
  trControl = ctrl
)

# Step 3: Train models using PCA data
set.seed(141)
model.results.pca = list()

model.results.pca[["Lasso"]] = train(
  x = x.pca, y = y.pca,
  method = "glmnet",
  trControl = ctrl,
  tuneLength = 10
)

model.results.pca[["SVM"]] = train(
  x = x.pca, y = y.pca,
  method = "svmLinear",
  trControl = ctrl
)

model.results.pca[["Decision Tree"]] = train(
  x = x.pca, y = y.pca,
  method = "rpart",
  trControl = ctrl
)

model.results.pca[["Random Forest"]] = train(
  x = x.pca, y = y.pca,
  method = "rf",
  trControl = ctrl
)

model.results.pca[["KNN"]] = train(
  x = x.pca, y = y.pca,
  method = "knn",
  trControl = ctrl,

```

```

tuneLength = 5
)

# Step 4: Compare PCA-based model accuracy
model.accuracy.pca = sapply(model.results.pca,
                             function(m) max(m$results$Accuracy))
model.accuracy.pca


##           Lasso           SVM Decision Tree Random Forest           KNN
## 0.9747373 0.8178947 0.7082610 0.9444737 0.7930206

# Step 5: Plot
# Extract cross-validated accuracy from each caret model
model.accuracy.pca <- sapply(model.results.pca, function(m) {
  if ("Accuracy" %in% colnames(m$results)) {
    max(m$results$Accuracy)
  } else {
    NA # Handle models like SVM (if probs or tuning failed)
  }
})

# Create data frame
accuracy_df <- data.frame(
  Model = names(model.accuracy.pca),
  Accuracy = as.numeric(model.accuracy.pca)
)

ggplot(accuracy_df, aes(x = reorder(Model, -Accuracy), y = Accuracy)) +
  geom_bar(stat = "identity", fill = "steelblue") +
  geom_text(aes(label = round(Accuracy, 3)), vjust = -0.5, size = 4) +
  ylim(0, 1) +
  theme_minimal() +
  labs(
    title = "Model Accuracy on PCA-Transformed Data",
    x = "Model",
    y = "Cross-Validated Accuracy"
  ) +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))

```



F1 & F2 & ROC Curve/AUC Score

Random Forest

```
# Main Results (Model Evaluation: Random Forest)
# Step 1: Store predicted labels and true test labels
true.label = factor(test.data$type)
rf.prediction = predict(model.results.pca[["Random Forest"]], pca.test)

# Step 2: Create dataframe for evaluation
evaluate.data = data.frame(
  truth = true.label,
  prediction = factor(rf.prediction, levels = levels(true.label))
)

# Step 3: F1
f1.macro = f_meas(evaluate.data, truth = truth,
  estimate = prediction, beta = 1)

# Step 4: F2
f2.macro = f_meas(evaluate.data, truth = truth,
  estimate = prediction, beta = 2)

# Step 5: Probabilities for AUC
rf.probs = predict(model.results.pca[["Random Forest"]], pca.test, type = "prob")
roc.multiclass = multiclass.roc(response = true.label,
  predictor = as.matrix(rf.probs))
```

```

auc.value = auc(roc.multiclass)

# Step 5b: Compute per-class AUCs
class_levels <- colnames(rf.probs)

auc_values <- sapply(class_levels, function(cls) {
  # Build a binary response: 1 for this class, 0 otherwise
  bin_truth <- as.numeric(true.label == cls)
  # Compute ROC and then its AUC
  roc_obj <- roc(bin_truth, rf.probs[[cls]], quiet = TRUE)
  auc(roc_obj)
})

# Make a little table
auc_table <- data.frame(
  Class = class_levels,
  AUC = round(auc_values, 3)
)
print(auc_table)

##           Class  AUC
## basal      basal 1.000
## cell_line cell_line 1.000
## HER          HER 1.000
## luminal_A luminal_A 0.991
## luminal_B luminal_B 0.985
## normal      normal 1.000

# Step 6: Output
f1.macro

## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 f_meas macro         0.908

f2.macro

## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 f_meas macro         0.910

auc.value

## Multi-class area under the curve: 0.9967

# Step 7: ROC Curve Plot
class_levels = colnames(rf.probs)

roc_df <- do.call(rbind, lapply(class_levels, function(class) {
  binary_response <- as.numeric(true.label == class)
  # Only compute ROC if both classes are present
  if (length(unique(binary_response)) < 2) {
    return(NULL) # Skip this class
  }

```

```

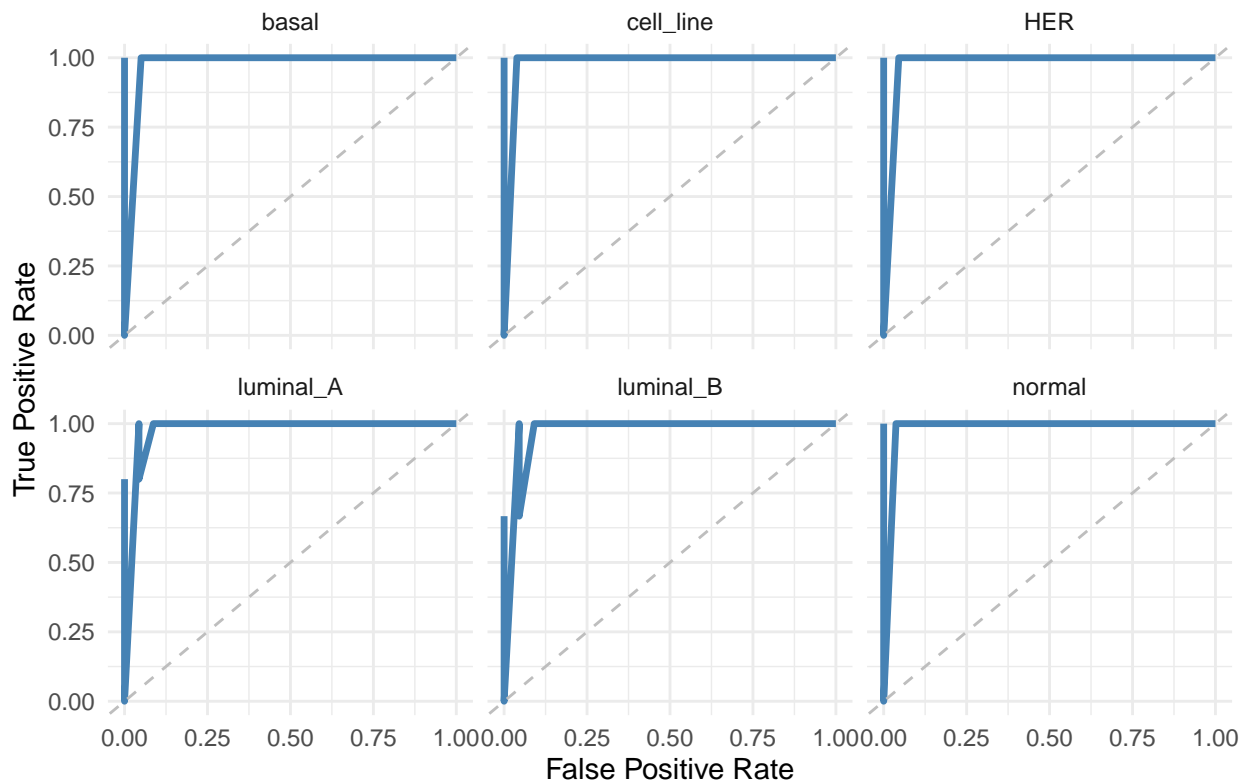
roc_obj <- roc(binary_response, rf.probs[[class]], quiet = TRUE)

data.frame(
  fpr = 1 - roc_obj$specificities,
  tpr = roc_obj$sensitivities,
  class = class,
  auc = rep(auc(roc_obj), length(roc_obj$sensitivities))
)
}))

ggplot(roc_df, aes(x = fpr, y = tpr)) +
  geom_line(linewidth = 1.2, color = "steelblue") +
  geom_abline(linetype = "dashed", color = "gray") +
  facet_wrap(~ class, ncol = 3) +
  theme_minimal() +
  labs(
    title = "Test Set: One-vs-Rest ROC Curves for Each Class (Random Forest)",
    x = "False Positive Rate",
    y = "True Positive Rate"
  )
)

```

Test Set: One-vs-Rest ROC Curves for Each Class (Random Forest)



Lasso

```

# Main Results (Model Evaluation: Lasso)
# Step 1: Store predicted labels and true test labels
true.label = factor(test.data$type)

```



```

lasso.prediction = predict(model.results.pca[["Lasso"]], pca.test)

# Step 2: Create dataframe for evaluation
evaluate.data = data.frame(
  truth = true.label,
  prediction = factor(lasso.prediction, levels = levels(true.label))
)

# Step 3: F1
f1.macro = f_meas(evaluate.data, truth = truth,
  estimate = prediction, beta = 1)

# Step 4: F2
f2.macro = f_meas(evaluate.data, truth = truth,
  estimate = prediction, beta = 2)

# Step 5: Probabilities for AUC
lasso.probs = predict(model.results.pca[["Lasso"]], pca.test, type = "prob")
roc.multiclass = multiclass.roc(response = true.label,
  predictor = as.matrix(lasso.probs))
auc.value = auc(roc.multiclass)

# After Step 5, you have:
lasso.probs <- predict(model.results.pca[["Lasso"]], pca.test, type = "prob")
true.label <- factor(test.data$type)

# Compute per-class AUCs:
class_levels <- colnames(lasso.probs)

auc_values <- sapply(class_levels, function(cls) {
  bin_truth <- as.numeric(true.label == cls)
  roc_obj <- roc(bin_truth, lasso.probs[[cls]], quiet = TRUE)
  auc(roc_obj)
})

# Build and print a table of AUCs
auc_table <- data.frame(
  Class = class_levels,
  AUC = round(auc_values, 3)
)
print(auc_table)

```

```

##           Class AUC
## basal      basal   1
## cell_line  cell_line 1
## HER        HER     1
## luminal_A  luminal_A 1
## luminal_B  luminal_B 1
## normal     normal   1

```

```

# Step 6: Output

```

```

f1.macro

```

```

## # A tibble: 1 x 3
##   .metric .estimator .estimate

```

```
##   <chr>   <chr>           <dbl>
## 1 f_meas macro           0.970
f2.macro

## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>           <dbl>
## 1 f_meas macro           0.971
auc.value

## Multi-class area under the curve: 1
# Step 7: ROC Curve Plot
class_levels = colnames(lasso.probs)

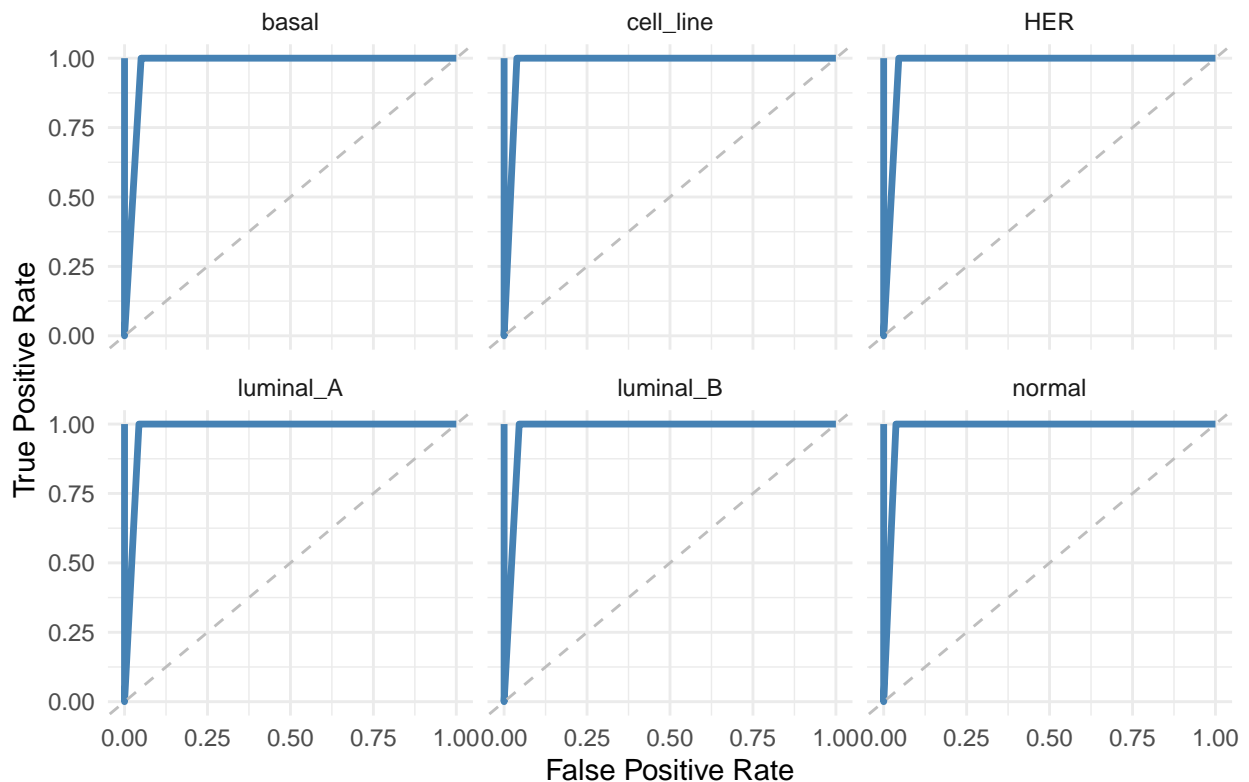
roc_df <- do.call(rbind, lapply(class_levels, function(class) {
  binary_response <- as.numeric(true.label == class)
  # Only compute ROC if both classes are present
  if (length(unique(binary_response)) < 2) {
    return(NULL) # Skip this class
  }

  roc_obj <- roc(binary_response, lasso.probs[[class]], quiet = TRUE)

  data.frame(
    fpr = 1 - roc_obj$specificities,
    tpr = roc_obj$sensitivities,
    class = class,
    auc = rep(auc(roc_obj), length(roc_obj$sensitivities))
  )
}))

ggplot(roc_df, aes(x = fpr, y = tpr)) +
  geom_line(linewidth = 1.2, color = "steelblue") +
  geom_abline(linetype = "dashed", color = "gray") +
  facet_wrap(~ class, ncol = 3) +
  theme_minimal() +
  labs(
    title = "Test Set: One-vs-Rest ROC Curves for Each Class (Lasso)",
    x = "False Positive Rate",
    y = "True Positive Rate"
  )
)
```

Test Set: One-vs-Rest ROC Curves for Each Class (Lasso)



Decision Tree

```
# Main Results (Model Evaluation: Decision Tree)
# Step 1: Store predicted labels and true test labels
true.label = factor(test.data$type)
tree.prediction = predict(model.results.pca[["Decision Tree"]], pca.test)

# Step 2: Create dataframe for evaluation
evaluate.data = data.frame(
  truth = true.label,
  prediction = factor(tree.prediction, levels = levels(true.label))
)

# Step 3: F1
f1.macro = f_meas(evaluate.data, truth = truth,
  estimate = prediction, beta = 1)

## Warning: While computing multiclass `precision()`, some levels had no predicted events
## (i.e. `true_positive + false_positive = 0`).
## Precision is undefined in this case, and those levels will be removed from the
## averaged result.
## Note that the following number of true events actually occurred for each
## problematic event level:
## 'HER': 6, 'luminal_B': 6

# Step 4: F2
f2.macro = f_meas(evaluate.data, truth = truth,
```

```

estimate = prediction, beta = 2)

## Warning: While computing multiclass `precision()`, some levels had no predicted events
## (i.e. `true_positive + false_positive = 0`).
## Precision is undefined in this case, and those levels will be removed from the
## averaged result.
## Note that the following number of true events actually occurred for each
## problematic event level:
## 'HER': 6, 'luminal_B': 6

# Step 5: Probabilities for AUC
tree.probs = predict(model.results.pca[["Decision Tree"]], pca.test, type = "prob")
roc.multiclass = multiclass.roc(response = true.label,
                                predictor = as.matrix(tree.probs))
auc.value = auc(roc.multiclass)

# Step 6: Output
f1.macro

## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 f_meas macro         0.838
f2.macro

## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>       <dbl>
## 1 f_meas macro         0.919
auc.value

## Multi-class area under the curve: 0.9333

# Step 7: ROC Curve Plot
class_levels = colnames(tree.probs)

roc_df <- do.call(rbind, lapply(class_levels, function(class) {
  binary_response <- as.numeric(true.label == class)
  # Only compute ROC if both classes are present
  if (length(unique(binary_response)) < 2) {
    return(NULL) # Skip this class
  }


  roc_obj <- roc(binary_response, tree.probs[[class]], quiet = TRUE)

  data.frame(
    fpr = 1 - roc_obj$specificities,
    tpr = roc_obj$sensitivities,
    class = class,
    auc = rep(auc(roc_obj), length(roc_obj$sensitivities))
  )
}))

ggplot(roc_df, aes(x = fpr, y = tpr)) +
  geom_line(linewidth = 1.2, color = "steelblue") +

```

```
geom_abline(linetype = "dashed", color = "gray") +
facet_wrap(~ class, ncol = 3) +
theme_minimal() +
labs(
  title = "Test Set: One-vs-Rest ROC Curves for Each Class (Decision Tree)",
  x = "False Positive Rate",
  y = "True Positive Rate"
)
```



KNN

```
# Main Results (Model Evaluation: KNN)
# Step 1: Store predicted labels and true test labels
true.label = factor(test.data$type)
knn.prediction = predict(model.results.pca[["KNN"]], pca.test)

# Step 2: Create dataframe for evaluation
evaluate.data = data.frame(
  truth = true.label,
  prediction = factor(knn.prediction, levels = levels(true.label))
)

# Step 3: F1
f1.macro = f_meas(evaluate.data, truth = truth,
  estimate = prediction, beta = 1)
```

```

# Step 4: F2
f2.macro = f_meas(evaluate.data, truth = truth,
                  estimate = prediction, beta = 2)

# Step 5: Probabilities for AUC
knn.probs = predict(model.results.pca[["KNN"]], pca.test, type = "prob")
roc.multiclass = multiclass.roc(response = true.label,
                               predictor = as.matrix(knn.probs))
auc.value = auc(roc.multiclass)

# Step 6: Output
f1.macro

## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>      <dbl>
## 1 f_meas macro        0.725
f2.macro

## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>   <chr>      <dbl>
## 1 f_meas macro        0.732
auc.value

## Multi-class area under the curve: 0.9623

# Step 7: ROC Curve Plot
class_levels = colnames(knn.probs)

roc_df <- do.call(rbind, lapply(class_levels, function(class) {
  binary_response <- as.numeric(true.label == class)
  # Only compute ROC if both classes are present
  if (length(unique(binary_response)) < 2) {
    return(NULL) # Skip this class
  }

  roc_obj <- roc(binary_response, knn.probs[[class]], quiet = TRUE)

  data.frame(
    fpr = 1 - roc_obj$specificities,
    tpr = roc_obj$sensitivities,
    class = class,
    auc = rep(auc(roc_obj), length(roc_obj$sensitivities))
  )
}))

ggplot(roc_df, aes(x = fpr, y = tpr)) +
  geom_line(linewidth = 1.2, color = "steelblue") +
  geom_abline(linetype = "dashed", color = "gray") +
  facet_wrap(~ class, ncol = 3) +
  theme_minimal() +
  labs(
    title = "Test Set: One-vs-Rest ROC Curves for Each Class (KNN)",

```

```

x = "False Positive Rate",
y = "True Positive Rate"
)

```

Test Set: One-vs-Rest ROC Curves for Each Class (KNN)

